

Joint Center for Satellite Data Assimilation

Office Note (unassigned)

THIS IS AN UNREVIEWED MANUSCRIPT, PRIMARILY INTENDED FOR INFORMAL
EXCHANGE OF INFORMATION AMONG JCSDA RESEARCHERS

CRTM: Cloud and Aerosol Optical Property Lookup Table Interpolation Speedup for REL-1.2

Paul van Delst^a
JCSDA/EMC/SAIC

December, 2008

^apaul.vandelst@noaa.gov

Change History

Date	Author	Change
2008-12-04	P.van Delst	Initial release.

1 Introduction

A central design principle of the components of the CRTM Tangent-linear (TL) and Adjoint (AD) models is that the Forward model is always called first. Thus, in principle, forward model calculations are always available for the various TL and AD components.

The REL-1.1 CRTM AtmScatter modules obtained the interpolated cloud and aerosol optical properties from the CloudCoeff and AerosolCoeff lookup tables (LUTs) as shown in figure 1.1. For the tangent-linear and adjoint models (figures 1.1(b) and (c)), the necessary interpolation parameters such as the bracketing indices and interpolating polynomials were always recomputed.

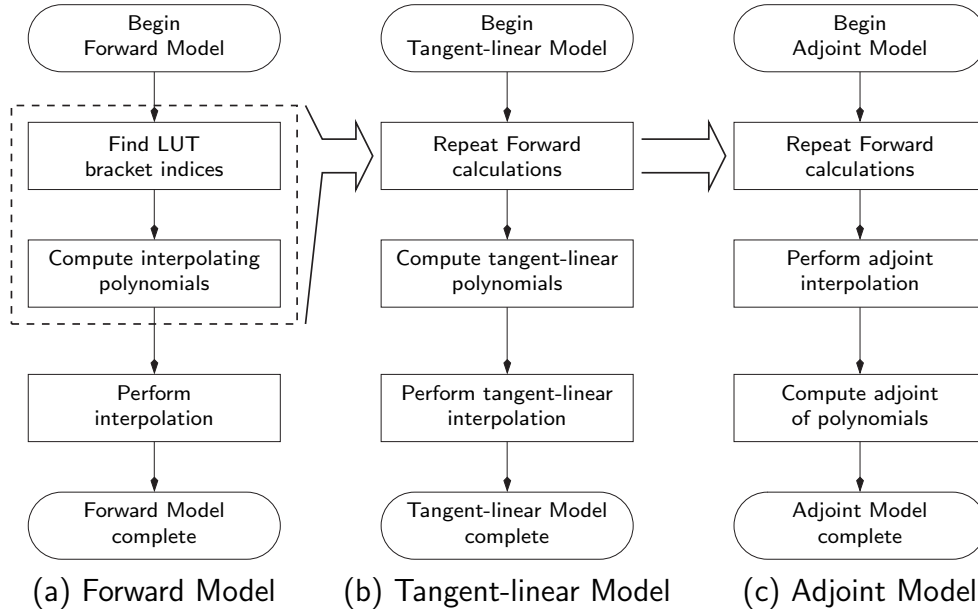


Figure 1.1: Schematic flowcharts of the REL-1.1 CRTM Forward, Tangent-linear, and Adjoint routines used to interpolate the cloud and aerosol optical properties from the CloudCoeff and AerosolCoeff lookup tables. The forward component of the interpolation is recomputed in the tangent-linear and adjoint routines.

The CRTM already saves the intermediate forward model results for its components in private structures, i.e. the structure definition is public, but the internals are private. We call these structures the internal variables for the particular CRTM component (e.g. AtmAbsorption, CloudScatter, AerosolScatter, etc).

To avoid the unnecessary forward model recalculations for the cloud and aerosol property LUT interpolations, separate structures were created to retain the intermediate LUT interpolation results, and added to the main internal variable structure. The structure used for cloudy atmospheres is shown in figure 1.2 and its usage in the CloudScatter internal variable structure is shown in figure 1.3. Similar, the aerosol structure is shown in figure 1.4, and its usage in the AerosolScatter internal variable structure is shown in figure 1.5. Because the cloud optical properties in the microwave are dependent upon an additional parameter (temperature), a different structure is defined for cloud and aerosol computations to minimise memory usage.

With the intermediate interpolation results now available for reuse, the tangent-linear and adjoint routines now avoid the forward model recalculations altogether, as shown schematically in figure 1.6.

The timing data that follows compares the CloudScatter and AerosolScatter component tests for a baseline (revision 2757) and modified (revision 2787) version of the CRTM library from the RB-1.2 branch.

```

TYPE :: CSinterp_type
  ! The interpolating polynomials
  TYPE(LPoly_type) :: wlp ! Frequency
  TYPE(LPoly_type) :: xlp ! Effective radius
  TYPE(LPoly_type) :: ylp ! Temperature
  ! The LUT interpolation indices
  INTEGER :: i1, i2      ! Frequency
  INTEGER :: j1, j2      ! Effective radius
  INTEGER :: k1, k2      ! Temperature
  ! The LUT interpolation boundary check
  LOGICAL :: f_outbound  ! Frequency
  LOGICAL :: r_outbound  ! Effective radius
  LOGICAL :: t_outbound  ! Temperature
  ! The interpolation input
  REAL(fp) :: f_int      ! Frequency
  REAL(fp) :: r_int      ! Effective radius
  REAL(fp) :: t_int      ! Temperature
  ! The data to be interpolated
  REAL(fp) :: f(NPTS)    ! Frequency
  REAL(fp) :: r(NPTS)    ! Effective radius
  REAL(fp) :: t(NPTS)    ! Temperature
END TYPE CSinterp_type

```

Figure 1.2: The structure definition to hold the forward model cloud optical properties lookup table interpolation results.

```

TYPE :: CRTM_CSVariables_type
  PRIVATE
  ! The interpolation data
  TYPE(CSinterp_type) :: csi(MAX_N_LAYERS, MAX_N_CLOUDS)
  ! The interpolation results
  REAL(fp) :: ke(MAX_N_LAYERS, MAX_N_CLOUDS) ! Mass extinction coefficient
  REAL(fp) :: w(MAX_N_LAYERS, MAX_N_CLOUDS) ! Single scatter albedo
  REAL(fp) :: g(MAX_N_LAYERS, MAX_N_CLOUDS) ! Asymmetry factor
  REAL(fp) :: pcoeff(0:MAX_N_LEGENDRE_TERMS,&
                     MAX_N_PHASE_ELEMENTS, &
                     MAX_N_LAYERS, &
                     MAX_N_CLOUDS) ! Phase coefficient
  ! The accumulated scattering coefficient
  REAL(fp) :: Total_bs(MAX_N_LAYERS) ! Volume scattering coefficient
END TYPE CRTM_CSVariables_type

```

Figure 1.3: The structure definition to hold all the forward model CloudScatter results showing the added csi structure array used to hold the intermediate interpolation results. The array dimensions are declared in the CRTM_Parameters module.

```

TYPE :: ASinterp_type
  ! The interpolating polynomials
  TYPE(LPoly_type) :: wlp ! Frequency
  TYPE(LPoly_type) :: xlp ! Effective radius
  ! The LUT interpolation indices
  INTEGER :: i1, i2      ! Frequency
  INTEGER :: j1, j2      ! Effective radius
  ! The LUT interpolation boundary check
  LOGICAL :: f_outbound  ! Frequency
  LOGICAL :: r_outbound  ! Effective radius
  ! The interpolation input
  REAL(fp) :: f_int      ! Frequency
  REAL(fp) :: r_int      ! Effective radius
  ! The data to be interpolated
  REAL(fp) :: f(NPTS)    ! Frequency
  REAL(fp) :: r(NPTS)    ! Effective radius
END TYPE ASinterp_type

```

Figure 1.4: The structure definition to hold the forward model aerosol optical properties lookup table interpolation results.

```

TYPE :: CRTM_ASVariables_type
  PRIVATE
  ! The interpolation data
  TYPE(ASinterp_type) :: asi(MAX_N_LAYERS, MAX_N_AEROSOLS)
  ! The interpolation result
  REAL(fp) :: ke(MAX_N_LAYERS, MAX_N_AEROSOLS) ! Mass extinction coefficient
  REAL(fp) :: w(MAX_N_LAYERS, MAX_N_AEROSOLS) ! Single scatter albedo
  REAL(fp) :: g(MAX_N_LAYERS, MAX_N_AEROSOLS) ! Asymmetry factor
  REAL(fp) :: pcoeff(0:MAX_N_LEGENDRE_TERMS,&
                     MAX_N_PHASE_ELEMENTS, &
                     MAX_N_LAYERS, &
                     MAX_N_AEROSOLS) ! Phase coefficient
  ! The accumulated scattering coefficient
  REAL(fp) :: Total_bs(MAX_N_LAYERS) ! Volume scattering coefficient
END TYPE CRTM_ASVariables_type

```

Figure 1.5: The structure definition to hold all the forward model AerosolScatter results showing the added asi structure array used to hold the intermediate interpolation results. The array dimensions are declared in the CRTM_Parameters module.

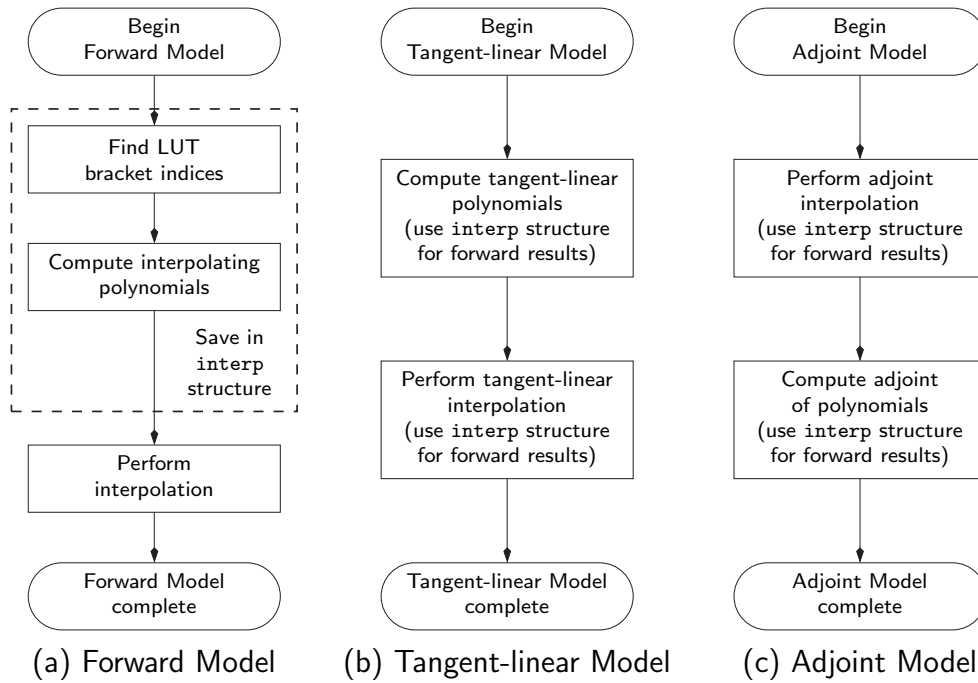


Figure 1.6: Schematic flowcharts of the REL-1.2 CRTM Forward, Tangent-linear, and Adjoint routines used to interpolate the cloud and aerosol optical properties from the CloudCoeff and AerosolCoeff lookup tables. The forward component of the interpolation is now saved in a structure variable and reused in the tangent-linear and adjoint routines.

2 CloudScatter Forward Model Profile Results

The forward model test results are shown separately for the microwave and infrared cloud optical property LUT interpolation. The microwave interpolation is done at different dimensionalities depending on the cloud type. The infrared interpolation is always two-dimensional.

No significant timing differences are seen in the forward model test since the number of calculations is the same, just that some of them are being retained in the `csi` structure as shown in figure 1.3.

2.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_mw	76.96	578.58	53.13	571.72
interp_3d	26.10	214.69	26.90	210.44
find_random_index	136.65	0.00	141.96	0.00
lpoly	20.74	96.16	17.99	91.97
interp_2d	24.73	56.77	23.12	56.60
interp_1d	2.74	0.00	2.74	0.00

Table 2.1: gfortran CloudScatter Forward model profile results for the Get_Cloud_Opt_MW subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_ir	31.69	189.73	23.17	186.95
interp_2d	30.76	70.61	28.76	70.40
find_random_index	47.62	0.00	49.47	0.00
lpoly	7.23	33.51	6.27	32.05

Table 2.2: gfortran CloudScatter Forward model profile results for the Get_Cloud_Opt_IR subroutine. All times in seconds.

2.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_mw	463.88	4633.15	399.15	4664.71
interp_3d	412.34	2133.74	390.48	2148.70
interp_2d	383.67	498.79	375.41	512.72
lpoly	165.98	598.86	168.51	600.09
find_random_index	432.79	0.00	461.62	0.00
interp_1d	6.98	0.00	7.18	0.00

Table 2.3: IBM CloudScatter Forward model profile results for the Get_Cloud_Opt_MW subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_ir	186.83	1514.98	164.67	1533.38
interp_2d	477.22	620.41	466.94	637.73
lpoly	57.84	208.69	58.72	209.12
find_random_index	150.82	0.00	160.87	0.00

Table 2.4: IBM CloudScatter Forward model profile results for the Get.Cloud.Opt.IR subroutine. All times in seconds.

3 CloudScatter Tangent-Linear Model Profile Results

As with the forward model, the tangent-linear test results are shown separately for the microwave and infrared cloud optical property LUT interpolation.

Here we begin to see significant differences in the routine timings between the Baseline and Modified test runs. Because the forward model interpolation parameters are reused rather than recalculated in the Modified runs, the time required for the `find_random_index` and `lpoly` routines do not contribute. Looking at the “self” value for the main routines, we see the Modified times are about 3x faster.

3.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_mw_tl</code>	305.41	1483.90	72.45	1280.26
<code>interp_3d_tl</code>	60.03	754.44	58.53	794.87
<code>interp_2d_tl</code>	53.45	184.56	53.27	199.80
<code>lpoly_tl</code>	34.28	149.79	31.44	142.35
<code>find_random_index</code>	140.22	0.00	-	-
<code>lpoly</code>	19.41	87.73	-	-

Table 3.1: gfortran CloudScatter Tangent-linear model profile results for the `Get_Cloud_Opt_MW_TL` subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_ir_tl</code>	117.27	447.04	27.28	375.34
<code>interp_2d_tl</code>	66.63	230.07	66.26	248.52
<code>lpoly_tl</code>	11.94	52.20	10.96	49.61
<code>find_random_index</code>	48.87	0.00	-	-
<code>lpoly</code>	6.76	30.57	-	-

Table 3.2: gfortran CloudScatter Tangent-linear model profile results for the `Get_Cloud_Opt_IR_TL` subroutine. All times in seconds.

3.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_mw_tl</code>	3886.24	15048.05	859.30	13631.37
<code>interp_3d_tl</code>	1031.52	8794.58	1026.03	8634.92
<code>interp_2d_tl</code>	950.29	1838.77	899.37	1836.93
<code>lpoly_tl</code>	193.76	1038.94	197.15	1036.98
<code>lpoly</code>	164.58	603.85	-	-
<code>find_random_index</code>	431.76	0.00	-	-

Table 3.3: IBM CloudScatter Tangent-linear model profile results for the `Get_Cloud_Opt_MW_TL` subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_ir_tl	1466.88	4324.64	352.51	3833.52
interp_2d_tl	1184.62	2292.19	1118.65	2284.80
lpoly_tl	67.52	362.05	68.70	361.37
lpoly	57.35	210.43	-	-
find_random_index	150.46	0.00	-	-

Table 3.4: IBM CloudScatter Tangent-linear model profile results for the Get_Cloud_Opt_IR_TL subroutine. All times in seconds.

4 CloudScatter Adjoint Model Profile Results

As with the other models, the adjoint test results are shown separately for the microwave and infrared cloud optical property LUT interpolation.

Here again we see significant differences in the routine timings between the Baseline and Modified test runs. Because the forward model interpolation parameters are reused rather than recalculated in the Modified runs, the time required for the `find_random_index` and `lpoly` routines do not contribute. Looking at the “self” value for the main routines, we see the Modified times are about 3-4x faster. A word of caution though: the test duration itself is quite short so the speedup factor should be regarded as a very rough guide. The main point is that the modifications do not make the code any slower.

4.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_mw_ad</code>	36.90	187.22	10.22	187.56
<code>interp_3d_ad</code>	8.70	98.79	8.79	120.22
<code>interp_2d_ad</code>	7.82	23.92	8.81	29.42
<code>find_random_index</code>	16.61	0.00	-	-
<code>lpoly_ad</code>	2.72	13.54	2.87	15.18
<code>lpoly</code>	2.25	10.33	-	-

Table 4.1: gfortran CloudScatter Adjoint model profile results for the `Get_Cloud_Opt_MW_AD` subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_ir_ad</code>	16.04	60.53	4.31	59.33
<code>interp_2d_ad</code>	9.74	29.82	10.96	36.60
<code>lpoly_ad</code>	1.66	8.25	1.75	9.24
<code>find_random_index</code>	5.79	0.00	-	-
<code>lpoly</code>	0.79	3.60	-	-

Table 4.2: gfortran CloudScatter Adjoint model profile results for the `Get_Cloud_Opt_IR_AD` subroutine. All times in seconds.

4.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_cloud_opt_mw_ad</code>	509.93	1842.51	101.19	1708.72
<code>interp_3d_ad</code>	126.71	1091.69	119.90	1100.07
<code>interp_2d_ad</code>	112.66	239.15	116.78	239.04
<code>lpoly_ad</code>	13.81	104.49	13.17	106.08
<code>lpoly</code>	19.90	69.37	-	-
<code>find_random_index</code>	51.25	0.00	-	-

Table 4.3: IBM CloudScatter Adjoint model profile results for the `Get_Cloud_Opt_MW_AD` subroutine. All times in seconds.

Routine	Baseline		Modified	
	self	called	self	called
get_cloud_opt_ir_ad	137.81	564.33	40.58	519.93
interp_2d_ad	140.44	298.13	145.25	297.33
lpoly_ad	8.42	63.68	8.01	64.56
lpoly	6.94	24.17	-	-
find_random_index	17.86	0.00	-	-

Table 4.4: IBM CloudScatter Adjoint model profile results for the Get_Cloud_Opt_IR_AD subroutine. All times in seconds.

5 AerosolScatter Forward Model Profile Results

Aerosol optical properties are only used in the infrared spectral region, so only one routine is profiled here – the generic `get_aerosol_opt`. As with the CloudScatter forward model results, no significant timing differences are seen since the number of calculations is essentially the same.

5.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_aerosol_opt</code>	39.52	230.06	34.28	245.13
<code>interp_2d</code>	40.73	91.91	37.85	103.73
<code>find_random_index</code>	49.10	0.00	53.28	0.00
<code>lpoly</code>	8.01	40.15	9.09	40.75
<code>aerosol_type_index</code>	0.18	0.00	0.44	0.00

Table 5.1: gfortran AerosolScatter Forward model profile results for the `Get_Aerosol_Opt` subroutine. All times in seconds.

5.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
<code>get_aerosol_opt</code>	267.45	1855.09	234.69	1889.82
<code>interp_2d</code>	579.35	817.80	583.05	795.91
<code>lpoly</code>	74.27	260.30	73.78	263.76
<code>find_random_index</code>	122.63	0.00	172.29	0.00
<code>aerosol_type_index</code>	0.74	0.00	1.03	0.00

Table 5.2: IBM AerosolScatter Forward model profile results for the `Get_Aerosol_Opt` subroutine. All times in seconds.

6 AerosolScatter Tangent-Linear Model Profile Results

Here we see the decreased execution time of the modified code due to the reuse of the intermediate interpolation results in the tangent-linear model. The speedup is around 3-4x.

6.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
get_aerosol_opt_tl	152.97	549.00	41.19	498.31
interp_2d_tl	85.45	285.46	98.44	315.26
lpoly_tl	15.46	66.40	15.01	69.23
find_random_index	51.30	0.00	-	-
lpoly	8.45	36.11	-	-
aerosol_type_index	0.38	0.00	0.37	0.00

Table 6.1: gfortran AerosolScatter Tangent-linear model profile results for the Get_Aerosol_Opt_TL subroutine. All times in seconds.

6.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
get_aerosol_opt_tl	1944.73	5353.09	436.08	4861.03
interp_2d_tl	1380.02	2944.39	1423.29	2896.04
lpoly_tl	85.19	480.71	86.00	454.49
lpoly	73.51	264.22	-	-
find_random_index	124.16	0.00	-	-
aerosol_type_index	0.88	0.00	1.21	0.00

Table 6.2: IBM AerosolScatter Tangent-linear model profile results for the Get_Aerosol_Opt_TL subroutine. All times in seconds.

7 AerosolScatter Adjoint Model Profile Results

As with the AerosolScatter tangent-linear model, we see the decreased execution time of the modified code due to the reuse of the intermediate interpolation results in the adjoint model. The speedup is around 2-3x, although the total test time is relatively short so care must be taken in interpreting those factors. As with the CloudScatter updates, the main point here is that the modifications do not make the code slower.

7.1 Linux test

Routine	Baseline		Modified	
	self	called	self	called
get_aerosol_opt_ad	28.84	117.95	9.86	92.47
interp_2d_ad	19.48	57.33	19.22	53.72
lpoly_ad	3.26	16.92	3.39	14.75
lpoly	1.80	7.77	-	-
find_random_index	9.38	0.00	-	-
aerosol_type_index	0.05	0.00	0.05	0.00

Table 7.1: gfortran AerosolScatter Adjoint model profile results for the Get_Aerosol_Opt_AD sub-routine. All times in seconds.

7.2 IBM test

Routine	Baseline		Modified	
	self	called	self	called
get_aerosol_opt_ad	290.81	1064.11	88.80	985.50
interp_2d_ad	269.14	561.43	276.07	563.02
lpoly_ad	15.95	124.51	15.52	121.79
lpoly	13.22	48.02	-	-
find_random_index	22.79	0.00	-	-
aerosol_type_index	0.19	0.00	0.16	0.00

Table 7.2: IBM AerosolScatter Adjoint model profile results for the Get_Aerosol_Opt_AD sub-routine. All times in seconds.